

Question 1: Do the following MongoDB operations.

1. Create a database "Student" with the following attributes Rollno, Age, ContactNo, Email-Id.

```
use collegedb;  
db.createCollection("Student");
```

2. Insert appropriate values

```
db.insert({_id:1, Rollno:1, Age:22, Name:"Prateek", ContactNo:991231233, EmailId :  
"prateek@gmail.com"});  
db.insert({_id:2, Rollno:10, Age:22, Name:"Mayank", ContactNo: 991231234,  
EmailId:"mayank@gmail.com"});  
db.insert({_id:3, Rollno:11, Age:22, Name:"ABC", ContactNo: 991231235,  
EmailId:"abc@gmail.com"});  
db.Sprateek@gmail.comtudent.find().pretty();
```

```
> db.createCollection("Student");  
{ "ok" : 1 }  
> db.Student.insert({"_id":1,"roll":1,"name":"prateek","age":21,"email":"prateek  
@gmail.com"});  
WriteResult({ "nInserted" : 1 })  
> db.Student.insert({"_id":2,"roll":2,"name":"rahul;", "age":21,"email":"rahul@gm  
ail.com"});  
WriteResult({ "nInserted" : 1 })  
> db.Student.insert({"_id":3,"roll":10,"name":"saif;", "age":21,"email":"saif@gma  
il.com"});  
WriteResult({ "nInserted" : 1 })  
> db.Student.insert({"_id":4,"roll":14,"name":"wali;", "age":21,"email":"wali@gma  
il.com"});  
WriteResult({ "nInserted" : 1 })
```

```
> db.Student.find()  
{ "_id" : 1, "roll" : 1, "name" : "prateek", "age" : 21, "email" : "prateek@gm  
ail.com" }  
{ "_id" : 2, "roll" : 2, "name" : "rahul;", "age" : 21, "email" : "rahul@gmail  
.com" }  
{ "_id" : 3, "roll" : 10, "name" : "saif;", "age" : 21, "email" : "saif@gmail.  
com" }  
{ "_id" : 4, "roll" : 14, "name" : "wali;", "age" : 21, "email" : "wali@gmail.  
com" }
```

3. Write a query to update the Email-Id of a student with roll 10.

```
db.Student.update({Rollno:10},{ $set:{EmailId:"mayankwali@gmail.com"}});
```

```
> db.Student.update({"roll":10},{ $set:{ "email": "saifnew@gmail.com" }});  
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })  
> db.Student.find({"roll":10})  
{ "_id" : 3, "roll" : 10, "name" : "saif;", "age" : 21, "email" : "saifnew@gmail.com" }  
>
```

4. Replace the student name from "ABC" to "FEM" of rollno 1.

```
db.Student.update({Rollno:11},{ $set:{Name:"FEM"}})
```

```
> db.Student.update({"name":"wali;"},{ $set:{ "name": "walinew" }});  
WriteResult({ "nMatched" : 1, "nUpserted" : 0, "nModified" : 1 })
```

5. Export the created table into the local file system.

```
mongoexport -d collegedb -c Student -f Rollno, Age, Name, EmailId --type=csv -o student.csv
```

```
prateek@ubuntu:~$ mongoexport -d collegedb -c Student -f name,age,contactNO,emailId --type=csv -o student.csv  
2020-10-08T15:08:43.450+0530    connected to: localhost  
2020-10-08T15:08:43.451+0530    exported 4 records  
prateek@ubuntu:~$ cat student.csv  
name,age,contactNO,emailId  
prateek,21,,  
rahul;,21,,  
saif;,21,,  
walinew,21,,
```

6. Drop the table

```
mongoexport -d collegedb -c Student -f Rollno, Age, Name, EmailId --type=csv -o student.csv
```

```
> show dbs  
admin          0.000GB  
collegedb      0.000GB  
company        0.000GB  
config         0.000GB  
local          0.000GB  
mydb           0.000GB  
> use collegedb  
switched to db collegedb  
> db.Student.drop()  
true
```

7. Import a given CSV dataset from the local file system into MongoDB collection.

mongoimport -d collegedb -c Student --type csv --file student.csv --headerline

```
prateek@ubuntu:~$ mongoimport -d collegedb -c Student --type csv --file student.csv --headerline
2020-10-08T16:01:52.805+0530    connected to: localhost
2020-10-08T16:01:53.047+0530    imported 4 documents
```

2. Perform the following DB operations using MongoDB.

1. Create a collection by name Customers with the following attributes.

Cust_id, Acc_Bal, Acc_Type

db.createCollection("Customers");

```
> db.createCollection("Customers");
{ "ok" : 1 }
```

2. Insert at least 5 values into the table

```
db.Customers.insert({_id:1,
custid:1,
accountbalance:2000,
accounttype:"closing"
});
db.Customers.insert({_id:2,
custid:2,
accountbalance:3000,
accounttype:"saving"
});
db.Customers.insert({_id:3,
custid:3,
accountbalance:4000,
accounttype:"closing"
});
db.Customers.insert({_id:4,
custid:4,
accountbalance:5000,
accounttype:"saving"
});
db.Customers.insert({_id:5,
custid:5,
accountbalance:6000,
accounttype:"closing"
});
```

```
> db.Customers.insert({"_id":1,"custid":1,"accountbalance":2000,"acctype":"saving"});
WriteResult({ "nInserted" : 1 })
> db.Customers.insert({"_id":21,"custid":2,"accountbalance":400,"acctype":"saving"});
uncaught exception: SyntaxError: missing : after property id :
@(shell):1:26
> db.Customers.insert({"_id":2,"custid":2,"accountbalance":400,"acctype":"saving"});
WriteResult({ "nInserted" : 1 })
> db.Customers.insert({"_id":3,"custid":3,"accountbalance":5000,"acctype":"saving"});
WriteResult({ "nInserted" : 1 })
> db.Customers.insert({"_id":4,"custid":4,"accountbalance":50000,"acctype":"saving"});
WriteResult({ "nInserted" : 1 })
> db.Customers.insert({"_id":5,"custid":5,"accountbalance":500000,"acctype":"saving"});
WriteResult({ "nInserted" : 1 })
```

3. Write a query to display those records whose total account balance is greater than 1200 of account type 'Z' for each customer_id.

```
db.Customers.find({
  accountbalance:{$gte:1200},
  accounttype:"saving"
});
```

```
> db.Customers.find({"accountbalance":{$gte:50000},"acctype":"saving"});
{ "_id" : 4, "custid" : 4, "accountbalance" : 50000, "acctype" : "saving" }
{ "_id" : 5, "custid" : 5, "accountbalance" : 500000, "acctype" : "saving" }
```

4. Determine Minimum and Maximum account balance for each customer_id.

```
db.Customers.aggregate(
[
{
$group:
{
_id:$custid,
max_bal:{$max:$accountbalance},
min_bal:{$min:$accountbalance}
}
}
])
```



```
> db.Customers.aggregate(
...   [
...     {
...       $group:
...       {
...         "_id": "$custid",
...         "max_bal": { $max: "$accountbalance" },
...         "min_bal": { $min: "$accountbalance" }
...       }
...     ]
...   );
{ "_id" : 1, "max_bal" : 2000, "min_bal" : 2000 }
{ "_id" : 5, "max_bal" : 500000, "min_bal" : 90000 }
{ "_id" : 2, "max_bal" : 400, "min_bal" : 400 }
{ "_id" : 4, "max_bal" : 50000, "min_bal" : 50000 }
{ "_id" : 6, "max_bal" : 50000, "min_bal" : 50000 }
{ "_id" : 3, "max_bal" : 5000, "min_bal" : 5000 }
```

5. Export created collection into the local file system

mongoexport -d labtestdb -c Customers -f custid, accountbalance, accounttype --type=csv -o Customers.csv.

6. Drop the table

db.Customers.drop();

7. Import a given CSV dataset from the local file system into MongoDB collection.

mongoimport -d labtestdb -c Customers --type csv --file Customers.csv --headerline

Cassandra Programs-To Do

Question 3. Perform the following DB operations using Cassandra.

```
prateek@ubuntu:~$ nodetool status
Datacenter: datacenter1
=====
Status=Up/Down
||/ State=Normal/Leaving/Joining/Moving
-- Address      Load          Tokens         Owns        Host ID
   Rack
UN  127.0.0.1    355.9 KiB     256            ?           cf2208aa-5850-48f9-bfcb-aa19c4e
10a5a rack1
```

1. Create a keyspace by name Employee

CREATE KEYSPACE Employee WITH replication = {'class':'SimpleStrategy','replication_factor':3};

```
cqlsh> create keyspace Employee with replication={'class':'SimpleStrategy', 'r
eplication_factor': 1};
```

2. Create a column family by name Employee-Info with attributes

Emp_Id Primary Key, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name

CREATE COLUMNFAMILY employee_info(emp_id INT PRIMARY KEY,emp_name VARCHAR,designation VARCHAR,doj VARCHAR,dept_name VARCHAR,salary INT);

```
cqlsh:employee> create table employee_info(  
... emp_id int PRIMARY KEY ,  
... emp_name varchar,  
... dept_name varchar,  
... designation varchar,  
... doj varchar,  
... salary int);
```

3. Insert the values into the table in batch

BEGIN BATCH

INSERT INTO

employee_info(emp_id,dept_name,designation,doj,emp_name,salary)values(120,'Development','CTO','10/11/2015','Prateek Aryan',2000000);

INSERT INTO

employee_info(emp_id,dept_name,designation,doj,emp_name,salary)values(121,'HR','Employee','20/01/2011','Mayank Wali',1500000);

INSERT INTO

employee_info(emp_id,dept_name,designation,doj,emp_name,salary)values(122,'Maintainance','staff','10/07/2020','Saifur Rahman',50000);

INSERT INTO

employee_info(emp_id,dept_name,designation,doj,emp_name,salary)values(123,'IT','Assistant','25/07/2014','Parth Chatta',100000);

APPLY BATCH;

select * from employee_info;

```
cqlsh:employee> BEGIN BATCH  
... insert INTO employee_info  
... (emp_id,emp_name,dept_name, designation , doj , salary )  
... VALUES (  
... 2,'Saifur Rahman','TE','head','01-02-20',5000);  
... insert INTO employee_info  
... (emp_id,emp_name,dept_name, designation , doj , salary )  
... VALUES (  
... 3,'Mayank Wali','ME','Teacher','02-03-1299',10000);  
... APPLY BATCH ;  
cqlsh:employee> SELECT * FROM employee_info ;
```

emp_id	dept_name	designation	doj	emp_name	salary
1	CSE	Student	01-09-29	Prateek Aryan	500
2	TE	head	01-02-20	Saifur Rahman	5000
3	ME	Teacher	02-03-1299	Mayank Wali	10000

3. Update Employee name and Department of Emp-Id 121

```
UPDATE employee_info SET emp_name = 'Aakash',dept_name = 'IT' WHERE emp_id = 121;  
select * from employee_info;
```

```
cqlsh:employee> UPDATE  
... employee_info  
... SET  
... emp_name='Prateek Aryan',dept_name='IT'  
... WHERE  
... emp_id =1;  
cqlsh:employee> select* FROM employee_info;
```

emp_id	dept_name	designation	doj	emp_name	salary
1	IT	Student	01-09-29	Prateek Aryan	500
2	TE	head	01-02-20	Saifur Rahman	5000
3	ME	Teacher	02-03-1299	Mayank Wali	10000

4. Sort the details of Employee records based on salary

```
cqlsh:employee> select* FROM employee_info ORDER BY salary;  
InvalidRequest: Error from server: code=2200 [Invalid query] message="ORDER BY  
is only supported when the partition key is restricted by an EQ or an IN."
```

5. Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.

```
ALTER TABLE employee_info ADD Project VARCHAR;
```

```
cqlsh:employee> ALTER TABLE employee_info  
... ADD project set <varchar >;  
cqlsh:employee> select * FROM employee_info ;
```

emp_id	dept_name	designation	doj	emp_name	project	salary
1	IT	Student	01-09-29	Prateek Aryan	null	500
2	TE	head	01-02-20	Saifur Rahman	null	5000
3	ME	Teacher	02-03-1299	Mayank Wali	null	10000

6. Update the altered table to add project names.

```
UPDATE employee_info SET project='TIP' WHERE emp_id=120;  
UPDATE employee_info SET project='Sentiment Analysis' WHERE emp_id=121;  
UPDATE employee_info SET project='Facial recognition' WHERE emp_id=123;  
select * from employee_info;
```



```
cqlsh:employee> SELECT * FROM employee_info ;
```

emp_id	dept_name	designation	doj	emp_name	project	salary
120	null	null	null	null	{'Investor Platform', 'Research Tool'}	null
1	IT	Student	01-09-29	Prateek Aryan	null	500
2	TE	head	01-02-20	Saifur Rahman	null	5000
3	ME	Teacher	02-03-1299	Mayank Wali	null	10000

7. Create a TTL of 15 seconds to display the values of Employees.

INSERT INTO employee_info(emp_id, dept_name, designation, doj, emp_name, project, salary) values(124, 'PR', 'Senior Manager', '8/8/2020', 'Load balancing server', 'Abhi', 20000) USING TTL 15;
SELECT TTL(designation) FROM employee_info where emp_id=124;

```
cqlsh:employee> BEGIN BATCH
... INSERT INTO
... employee_info (emp_id, emp_name, designation, doj , salary , dept_name )
... VALUES ( 123, 'Harsh', 'Master', '05-09-23', 60000, 'ME') USING TTL 15;
... APPLY BATCH ;
cqlsh:employee> SELECT TTL(designation) FROM employee_info WHERE emp_id=123;
InvalidRequest: Error from server: code=2200 [Invalid query] message="Undefined column name emp_id"
cqlsh:employee> SELECT TTL(designation) FROM employee_info WHERE emp_id=123;

ttl(designation)
-----
(0 rows)
```

The screenshot shows the Visual Studio Code interface with a terminal window open. The terminal displays the following CQL queries and their results:

```
cqlsh> DESCRIBE KEYSPACES;
system_schema system test keyspace system_traces
system_auth library system distributed

cqlsh> CREATE KEYSPACE Employee WITH replication = {'class': 'SimpleStrategy', 'replication factor': 3};
InvalidRequest: Error from server: code=2200 [Invalid query] message="No keyspace has been specified. USE a keyspace, or explicitly specify keyspace.tablename"

cqlsh> use employee ;
cqlsh:employee> CREATE COLUMNFAMILY employee_info(emp_id INT PRIMARY KEY, emp_name VARCHAR, designation VARCHAR, doj VARCHAR, dept_name VARCHAR, salary INT);
cqlsh:employee> BEGIN BATCH
... INSERT INTO employee_info(emp_id, dept_name, designation, doj, emp_name, salary) values(120, 'Development', 'CTO', '18/11/2015', 'Prateek Aryan', 2000000);
... INSERT INTO employee_info(emp_id, dept_name, designation, doj, emp_name, salary) values(121, 'HR', 'Employee', '20/01/2011', 'Mayank Wali', 1500000);
... INSERT INTO employee_info(emp_id, dept_name, designation, doj, emp_name, salary) values(122, 'Maintenance', 'staff', '10/07/2020', 'Saifur Rahman', 50000);
... INSERT INTO employee_info(emp_id, dept_name, designation, doj, emp_name, salary) values(123, 'IT', 'Assistant', '25/07/2014', 'Parth Chatta', 100000);
... APPLY BATCH;
cqlsh:employee> select * from employee_info;

emp_id | dept_name | designation | doj | emp_name | salary
-----|-----|-----|-----|-----|-----
120 | Development | CTO | 18/11/2015 | Prateek Aryan | 2000000
123 | IT | Assistant | 25/07/2014 | Parth Chatta | 100000
122 | Maintenance | staff | 10/07/2020 | Saifur Rahman | 50000
121 | HR | Employee | 20/01/2011 | Mayank Wali | 1500000
(4 rows)

cqlsh:employee> UPDATE employee_info SET emp_name = 'Aakash', dept_name = 'IT' WHERE emp_id = 121;
cqlsh:employee> ALTER TABLE employee_info ADD Project VARCHAR;
cqlsh:employee> UPDATE employee_info SET project='TIP' WHERE emp_id=120;
cqlsh:employee> UPDATE employee_info SET project='Sentiment Analysis' WHERE emp_id=121;
cqlsh:employee> UPDATE employee_info SET project='Facial recognition' WHERE emp_id=123;
cqlsh:employee> select * from employee_info;

emp_id | dept_name | designation | doj | emp_name | project | salary
-----|-----|-----|-----|-----|-----|-----
120 | Development | CTO | 18/11/2015 | Prateek Aryan | TIP | 2000000
123 | IT | Assistant | 25/07/2014 | Parth Chatta | Facial recognition | 100000
122 | Maintenance | staff | 10/07/2020 | Saifur Rahman | null | 50000
121 | IT | Employee | 20/01/2011 | Aakash | Sentiment Analysis | 1500000
(4 rows)

cqlsh:employee> INSERT INTO employee_info(emp_id, dept_name, designation, doj, emp_name, project, salary) values(124, 'PR', 'Senior Manager', '8/8/2020', 'Load balancing server', 'Abhi', 20000) USING TTL 15;
cqlsh:employee> SELECT TTL(designation) FROM employee_info where emp_id=124;

ttl(designation)
-----
7
(1 rows)

cqlsh:employee>
```


Question 4. Perform the following DB operations using Cassandra.

1. Create a keyspace by name Library

```
CREATE KEYSPACE library WITH replication = { 'class':'SimpleStrategy','replication_factor':3};
```

2. Create a column family by name Library-Info with attributes

Stud_Id Primary Key,

Counter_value of type Counter,

Stud_Name, Book-Name, Book-Id, Date_of_issue

```
CREATE COLUMNFAMILY library_info(stud_id int, counter_value counter, stud_name text,  
book_name text, book_id int, date_of_issue text, PRIMARY KEY(stud_id, stud_name, book_name,  
book_id, date_of_issue));
```

```
cqlsh:employee> CREATE KEYSPACE library WITH replication = {'class': 'SimpleStrategy'  
, 'replication_factor': 3};  
cqlsh:employee> CREATE COLUMNFAMILY library_info(  
... stud_id uuid,  
... counter_value counter,  
... stud_name varchar ,  
... book_name varchar ,  
... book_id int,  
... doi varchar ,  
... primary KEY (stud_id, stud_name, book_name, book_id,doi));
```

3. Insert the values into the table in batch

```
UPDATE library_info SET counter_value = counter_value +1 WHERE stud_id = 111 and  
stud_name = 'Prateek' and book_name='Machine Learning' and book_id=4567 and  
date_of_issue = '20-10-2020';
```

```
UPDATE library_info SET counter_value = counter_value +1 WHERE stud_id = 112 and  
stud_name = 'Mayank Wali' and book_name='Social Studies' and book_id=4589 and  
date_of_issue = '17-09-2020';
```

```
UPDATE library_info SET counter_value = counter_value +1 WHERE stud_id = 125 and  
stud_name = 'Parth Chatta' and book_name='Mathematics' and book_id=4555 and  
date_of_issue = '31-11-2020';
```

```
UPDATE library_info SET counter_value = counter_value +1 WHERE stud_id = 126 and  
stud_name = 'Saifur Rahman' and book_name='Science' and book_id=4557 and date_of_issue =  
'31-12-2020';
```

```
cqlsh:library> UPDATE library_info SET counter_value = counter_value +1 WHERE
stud_id = 111 and stud_name = 'Prateek' and book_name='Machine Learning' and
book_id=4567 and date_of_issue = '20-10-2020';
cqlsh:library> UPDATE library_info SET counter_value = counter_value +1 WHERE
stud_id = 112 and stud_name = 'Mayank Wali' book_name='Social Studies' and boo
k_id=4589 and date_of_issue = '17-09-2020';
SyntaxException: line 1:107 mismatched input 'book_name' expecting EOF (...112
and stud_name = 'Mayank Wali' [book_name]...)
cqlsh:library> UPDATE library_info SET counter_value = counter_value +1 WHERE
stud_id = 112 and stud_name = 'Mayank Wali' and book_name='Social Studies' and
book_id=4589 and date_of_issue = '17-09-2020';
cqlsh:library> UPDATE library_info SET counter_value = counter_value +1 WHERE
stud_id = 125 and stud_name = 'Parth Chatta' and book_name='Mathematics' and b
ook_id=4555 and date_of_issue = '31-11-2020';
cqlsh:library> UPDATE library_info SET counter_value = counter_value +1 WHERE
stud_id = 126 and stud_name = 'Saifur Rahman' and book_name='Science' and book
_id=4557 and date_of_issue = '31-12-2020';
cqlsh:library> DESCRIBE library_info;
```

3. Display the details of the table created and increase the value of the counter

DESCRIBE library_info;

```
CREATE TABLE library.library_info (
  stud_id int,
  stud_name text,
  book_name text,
  book_id int,
  date_of_issue text,
  counter_value counter,
  PRIMARY KEY (stud_id, stud_name, book_name, book_id, date_of_issue)
) WITH CLUSTERING ORDER BY (stud_name ASC, book_name ASC, book_id ASC, date_of
_issue ASC)
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredC
ompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassan
dra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND dclocal_read_repair_chance = 0.1
AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair_chance = 0.0
AND speculative_retry = '99PERCENTILE';
```

4. Write a query to show that a student with id 112 has taken a book "BDA" 2 times.

```
SELECT * FROM library_info WHERE stud_id=112;
```

```
cqlsh:library> SELECT *  
... FROM library_info WHERE stud_id=112;  
  
stud_id | stud_name | book_name | book_id | date_of_issue | counter_value  
-----+-----+-----+-----+-----+-----  
112 | Mayank Wali | Social Studies | 4589 | 17-09-2020 | 1
```

5. Export the created column to a CSV file.

```
COPY library_info TO '/home/prateek/library.csv';
```

```
cqlsh:library> COPY library_info TO '/home/prateek/library.csv'  
... ;  
Using 3 child processes  
  
Starting copy of library.library_info with columns [stud_id, stud_name, book_name, book_id, date_of_issue, counter_value].  
Processed: 4 rows; Rate: 11 rows/s; Avg. rate: 11 rows/s  
4 rows exported to 1 files in 0.429 seconds.
```

6. Import a given csv dataset from local file system into Cassandra column family

```
CREATE COLUMNFAMILY library_info_duplicate(stud_id int, counter_value counter, stud_name  
text, book_name text, book_id int, date_of_issue text, PRIMARY KEY(stud_id, stud_name,  
book_name, book_id, date_of_issue));
```

```
COPY library_info_duplicate FROM '/home/prateek/library.csv';
```

```
SELECT * FROM library_info_duplicate;
```

```
cqlsh:library> CREATE COLUMNFAMILY library_info_duplicate(stud_id int, counter_value counter, stud_name text, book_name text, book_id int, date_of_issue text, PRIMARY KEY(stud_id, stud_name, book_name, book_id, date_of_issue));
cqlsh:library> COPY library_info_duplicate FROM '/home/prateek/library.csv';
Using 3 child processes
```

Starting copy of library.library_info_duplicate with columns [stud_id, stud_name, book_name, book_id, date_of_issue, counter_value].

Processed: 4 rows; Rate: 3 rows/s; Avg. rate: 5 rows/s

4 rows imported from 1 files in 0.810 seconds (0 skipped).

```
cqlsh:library> select * from library
```

```
library.          library_info          library_info_duplicate
```

```
cqlsh:library> select * from library_info_duplicate
```

```
... ;
```

stud_id	stud_name	book_name	book_id	date_of_issue	counter_value
---------	-----------	-----------	---------	---------------	---------------

125	Parth Chatta	Mathematics	4555	31-11-2020	1
-----	--------------	-------------	------	------------	---

111	Prateek	Machine Learning	4567	20-10-2020	1
-----	---------	------------------	------	------------	---

112	Mayank Wali	Social Studies	4589	17-09-2020	1
-----	-------------	----------------	------	------------	---

126	Saifur Rahman	Science	4557	31-12-2020	1
-----	---------------	---------	------	------------	---

(4 rows)

Question 5. Develop a MapReduce program to count the number of occurrences of words in a given file.

WCMapper.java

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reporter;

public class WCMapper extends MapReduceBase implements Mapper<LongWritable,
    Text, Text, IntWritable> {

    // Map function
    public void map(LongWritable key, Text value, OutputCollector<Text,
        IntWritable> output, Reporter rep) throws IOException
    {

        String line = value.toString();

        // Splitting the line on spaces
        for (String word : line.split(" "))
        {
            if (word.length() > 0)
            {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

WCReducer.java

```
import java.io.IOException;
import java.util.Iterator;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text,
    IntWritable, Text, IntWritable> {
```

```
// Reduce function
public void reduce(Text key, Iterator<IntWritable> value,
                  OutputCollector<Text, IntWritable> output,
                  Reporter rep) throws IOException
{
    int count = 0;

    // Counting the frequency of each words
    while (value.hasNext())
    {
        IntWritable i = value.next();
        count += i.get();
    }

    output.collect(key, new IntWritable(count));
}
}
```

WCDriver.java

```
import java.io.IOException;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WCDriver extends Configured implements Tool {

    public int run(String args[]) throws IOException
    {
        if (args.length < 2)
        {
            System.out.println("Please give valid inputs");
            return -1;
        }

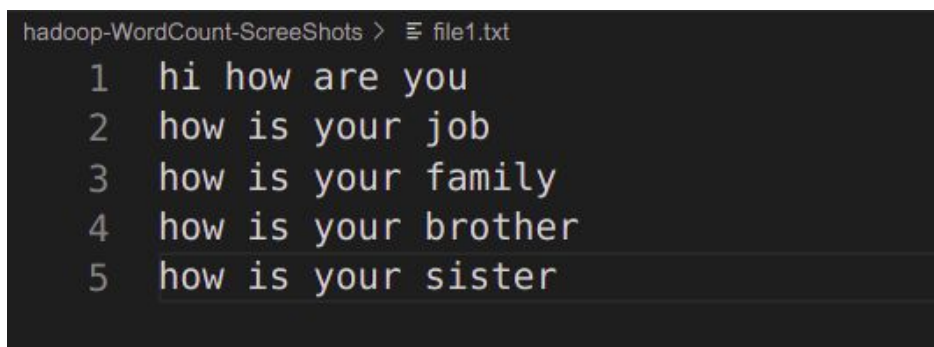
        JobConf conf = new JobConf(WCDriver.class);
        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new Path(args[1]));
    }
}
```

```
        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class);
        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.class);
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.class);
        JobClient.runJob(conf);
        return 0;
    }

    // Main Method
    public static void main(String args[]) throws Exception
    {
        int exitCode = ToolRunner.run(new WCDriver(), args);
        System.out.println(exitCode);
    }
}
```

HADOOP COMMANDS :

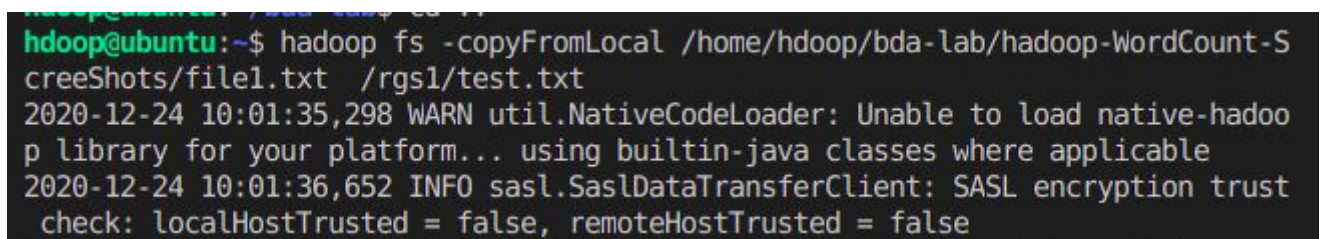
[file1.txt](#)



```
hadoop-WordCount-ScreenShots > file1.txt
1 hi how are you
2 how is your job
3 how is your family
4 how is your brother
5 how is your sister
```

[hadoop fs -copyFromLocal](#)

[/home/hadoop/bda-lab/hadoop-WordCount-ScreenShots/file1.txt /rgs1/test.txt](#)



```
hadoop@ubuntu:~$ hadoop fs -copyFromLocal /home/hadoop/bda-lab/hadoop-WordCount-S
creeShots/file1.txt /rgs1/test.txt
2020-12-24 10:01:35,298 WARN util.NativeCodeLoader: Unable to load native-hadoo
p library for your platform... using builtin-java classes where applicable
2020-12-24 10:01:36,652 INFO sasl.SaslDataTransferClient: SASL encryption trust
check: localhostTrusted = false, remoteHostTrusted = false
```

`hadoop jar /home/hadoop/bda-lab/hadoop-WordCount-ScreenShots/wordcount.jar WordCount /rgs1/test.txt /rgs1/output`

```
hadoop@ubuntu:~$ hadoop jar /home/hadoop/bda-lab/hadoop-WordCount-ScreenShots/wordcount.jar WordCount /rgs1/test.txt /rgs1/output
2020-12-24 10:17:06,637 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform..
. using builtin-java classes where applicable
2020-12-24 10:17:07,156 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2020-12-24 10:17:07,242 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2020-12-24 10:17:07,242 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2020-12-24 10:17:07,453 INFO input.FileInputFormat: Total input files to process : 1
2020-12-24 10:17:08,160 INFO mapreduce.JobSubmitter: number of splits:1
2020-12-24 10:17:08,763 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1012990088_0001
2020-12-24 10:17:08,763 INFO mapreduce.JobSubmitter: Executing with tokens: []
2020-12-24 10:17:08,977 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2020-12-24 10:17:08,978 INFO mapreduce.Job: Running job: job_local1012990088_0001
2020-12-24 10:17:09,013 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2020-12-24 10:17:09,026 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2020-12-24 10:17:09,026 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders
under output directory:false, ignore cleanup failures: false
2020-12-24 10:17:09,026 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output
t.FileOutputCommitter
2020-12-24 10:17:09,272 INFO mapred.LocalJobRunner: Waiting for map tasks
2020-12-24 10:17:09,273 INFO mapred.LocalJobRunner: Starting task: attempt_local1012990088_0001_m_000000_0
2020-12-24 10:17:09,334 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2020-12-24 10:17:09,334 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders
under output directory:false, ignore cleanup failures: false
2020-12-24 10:17:09,423 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2020-12-24 10:17:09,426 INFO mapred.MapTask: Processing split: hdfs://127.0.0.1:9000/rgs1/test.txt:0+88
2020-12-24 10:17:09,601 INFO mapred.MapTask: (EQUATOR) 0 kvi 26214396(104857584)
2020-12-24 10:17:09,601 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2020-12-24 10:17:09,601 INFO mapred.MapTask: soft limit at 83886080
2020-12-24 10:17:09,601 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2020-12-24 10:17:09,601 INFO mapred.MapTask: kvstart = 26214396; length = 6553600
2020-12-24 10:17:09,605 INFO mapred.MapTask: Map output collector class = org.apache.hadoop.mapred.MapTask$M
apOutputBuffer
2020-12-24 10:17:09,693 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = fa
lse, remoteHostTrusted = false
2020-12-24 10:17:10,001 INFO mapreduce.Job: Job job_local1012990088_0001 running in uber mode : false
2020-12-24 10:17:10,004 INFO mapreduce.Job: map 0% reduce 0%
2020-12-24 10:17:10,047 INFO mapred.LocalJobRunner:
2020-12-24 10:17:10,055 INFO mapred.MapTask: Starting flush of map output
2020-12-24 10:17:10,057 INFO mapred.MapTask: Spilling map output
2020-12-24 10:17:10,057 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufvoid = 104857600
2020-12-24 10:17:10,058 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214320(104857280); len
gth = 77/6553600
2020-12-24 10:17:10,144 INFO mapred.MapTask: Finished spill 0
2020-12-24 10:17:10,169 INFO mapred.Task: Task:attempt_local1012990088_0001_m_000000_0 is done. And is in th
e process of committing
2020-12-24 10:17:10,186 INFO mapred.LocalJobRunner: map
2020-12-24 10:17:10,186 INFO mapred.Task: Task 'attempt_local1012990088_0001_m_000000_0' done.
2020-12-24 10:17:10,198 INFO mapred.Task: Final Counters for attempt_local1012990088_0001_m_000000_0: Counte
```


hadoop fs -ls /rgs1/output/part-r-00000

```
hadoop@ubuntu:~$ hadoop fs -cat /rgs1/output/part-r-00000
2020-12-24 10:20:13,595 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform..
. using builtin-java classes where applicable
2020-12-24 10:20:14,272 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false,
remoteHostTrusted = false
are      1
brother  1
family   1
hi        1
how       5
is        4
job       1
sister   1
you       1
your      4
hadoop@ubuntu:~$
```

Question 6. For the given file, Create a Map Reduce program.

a) Find the average temperature for each year from the NCDC data set.

AverageMapper.java

```
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class AverageMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{

    public static final int MISSING = 9999;

    public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException
    {
        String line = value.toString();
        String year = line.substring(15,19);
        int temperature;
        if (line.charAt(87)=='+')
            temperature = Integer.parseInt(line.substring(88, 92));
        else
```

```
        temperature = Integer.parseInt(line.substring(87, 92));

        String quality = line.substring(92, 93);
        if(temperature != MISSING && quality.matches("[01459]"))
            context.write(new Text(year),new IntWritable(temperature));
    }
}
```

AverageReducer.java

```
import org.apache.hadoop.mapreduce.*;
import java.io.IOException;

public class AverageReducer extends Reducer <Text, IntWritable,Text, IntWritable >
{
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
        InterruptedException
    {
        int max_temp = 0;
        int count = 0;
        for (IntWritable value : values)
        {
            max_temp += value.get();
            count+=1;
        }
        context.write(key, new IntWritable(max_temp/count));
    }
}
```

AverageDriver.java

```
import org.apache.hadoop.io.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class AverageDriver
{

    public static void main (String[] args) throws Exception
    {
        if (args.length != 2)
        {
            System.err.println("Please Enter the input and output parameters");
            System.exit(-1);
        }
    }
}
```

```
        Job job = new Job();
        job.setJarByClass(AverageDriver.class);
        job.setJobName("Max temperature");

        FileInputFormat.addInputPath(job,new Path(args[0]));
        FileOutputFormat.setOutputPath(job,new Path (args[1]));

        job.setMapperClass(AverageMapper.class);
        job.setReducerClass(AverageReducer.class);

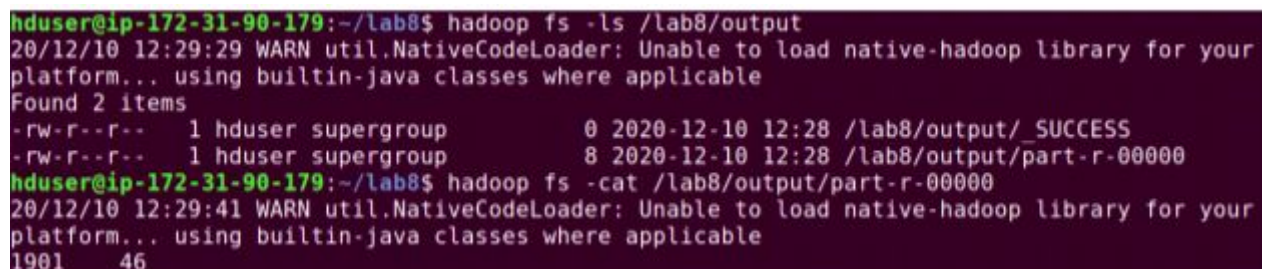
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        System.exit(job.waitForCompletion(true)?0:1);
    }
}
```

Hadoop Commands :

```
start-dfs.sh
start-yarn.sh
hdfs dfs -put /home/hdoop/Average/1909.txt
hadoop jar /home/hdoop/bda-lab/average/wordcount.jar AverageDriver /average/test.txt
/average/output
```

Creating a jar file.



```
hduser@ip-172-31-90-179:~/lab8$ hadoop fs -ls /lab8/output
20/12/10 12:29:29 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 hduser supergroup      0 2020-12-10 12:28 /lab8/output/_SUCCESS
-rw-r--r--  1 hduser supergroup    8 2020-12-10 12:28 /lab8/output/part-r-000000
hduser@ip-172-31-90-179:~/lab8$ hadoop fs -cat /lab8/output/part-r-000000
20/12/10 12:29:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
1901    46
```

```
HDFS: Number of bytes written=8
HDFS: Number of read operations=13
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
Map-Reduce Framework
  Map input records=6565
  Map output records=6564
  Map output bytes=59076
  Map output materialized bytes=72210
  Input split bytes=97
  Combine input records=0
  Combine output records=0
  Reduce input groups=1
  Reduce shuffle bytes=72210
  Reduce input records=6564
  Reduce output records=1
  Spilled Records=13128
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=53
  CPU time spent (ms)=0
  Physical memory (bytes) snapshot=0
  Virtual memory (bytes) snapshot=0
  Total committed heap usage (bytes)=242360320
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=888190
File Output Format Counters
  Bytes Written=8
hduser@ip-172-31-90-179:~/lab8$
```

```
hduser@ip-172-31-90-179:~/lab8$ hadoop fs -ls /lab8/output
20/12/10 12:29:29 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r--  1 hduser supergroup          0 2020-12-10 12:28 /lab8/output/_SUCCESS
-rw-r--r--  1 hduser supergroup          8 2020-12-10 12:28 /lab8/output/part-r-000000
hduser@ip-172-31-90-179:~/lab8$ hadoop fs -cat /lab8/output/part-r-000000
20/12/10 12:29:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your
platform... using builtin-java classes where applicable
1901      46
```

OUTPUT :
1901 46

Question 7. Write Queries in Hive to do the following.

1. Create an external table named with the following attributes

-> Empl_ID ->Emp_Name -> Designation -> Salary

```
CREATE DATABASE IF NOT EXISTS lab9 COMMENT 'employee program' WITH DBPROPERTIES ('creator'=PRATEEK);  
SHOW DATABASES;  
DESCRIBE DATABASE lab9;  
USE lab9;  
CREATE EXTERNAL TABLE IF NOT EXISTS Employee(EmpID INT,EmpName STRING, Designation STRING,Salary FLOAT) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

2. Load data into table from a given file

```
LOAD DATA LOCAL INPATH '/home/prateek/Downloads/employeeInput.txt' OVERWRITE INTO TABLE Employee;  
SELECT * FROM Employee;
```

3. Create a view to Generate a query to retrieve the employee details who earn a salary of more than Rs 30000.

```
CREATE VIEW emp_30000 AS SELECT * FROM Employee WHERE Salary>30000;  
SELECT * FROM emp_30000;
```

4. Alter the table to add a column Dept_Id and Generate a query to retrieve the employee details in order by using Dept_Id

```
ALTER TABLE Employee ADD COLUMNS(DeptID INT);  
LOAD DATA LOCAL INPATH '/home/prateek/Downloads/employeeInputAltered.txt' OVERWRITE INTO TABLE Employee;  
SELECT * FROM Employee;  
SELECT * FROM Employee ORDER BY DeptID;
```

5. Generate a query to retrieve the number of employees in each department whose salary is greater than 30000

```
SELECT DeptID,count(*) FROM Employee WHERE Salary>=30000 GROUP BY DeptID;
```

6. Create another table Department with attributes

-> Dept_Id ->Dept_name ->Emp_Id

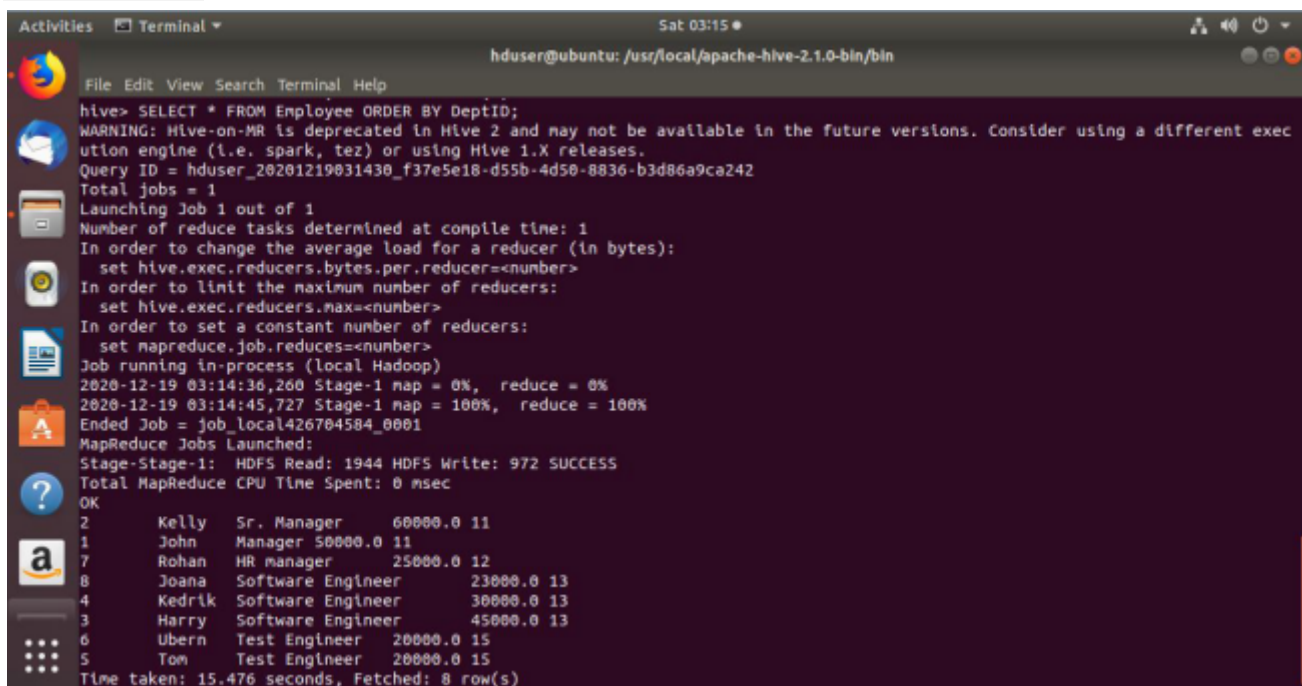
```
CREATE EXTERNAL TABLE IF NOT EXISTS Department(DeptId INT,DeptName STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

```
LOAD DATA LOCAL INPATH '/home/prateek/Downloads/DepartmentInput.txt' OVERWRITE  
INTO TABLE Department;  
SELECT * FROM Department;
```

7.Display the cumulative details of each employee along with department details

```
SELECT a.EmpID,a.EmpName,a.Designation,a.Salary,b.DeptName FROM Employee a  
JOIN Department b ON a.DeptID=b.DeptID;
```

Screenshots:



The screenshot shows a terminal window with the following output:

```
hive> SELECT * FROM Employee ORDER BY DeptID;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.  
Query ID = hduser_20201219031430_f37e5e18-d55b-4d50-8836-b3d86a9ca242  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks determined at compile time: 1  
In order to change the average load for a reducer (in bytes):  
  set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
  set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
  set mapreduce.job.reduces=<number>  
Job running in-process (local Hadoop)  
2020-12-19 03:14:36,260 Stage-1 map = 0%, reduce = 0%  
2020-12-19 03:14:45,727 Stage-1 map = 100%, reduce = 100%  
Ended Job = job_local426704584_0001  
MapReduce Jobs Launched:  
Stage-Stage-1:  HDFS Read: 1944 HDFS Write: 972 SUCCESS  
Total MapReduce CPU Time Spent: 0 msec  
OK  
2      Kelly      Sr. Manager      60000.0 11  
1      John       Manager 50000.0 11  
7      Rohan      HR manager       25000.0 12  
8      Joana      Software Engineer 23000.0 13  
4      Kedrik     Software Engineer 30000.0 13  
3      Harry      Software Engineer 45000.0 13  
6      Ubern      Test Engineer    20000.0 15  
5      Tom        Test Engineer    20000.0 15  
Time taken: 15.476 seconds, Fetched: 8 row(s)
```

```
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Job running in-process (local Hadoop)
2020-12-19 03:18:13,778 Stage-3 map = 100%, reduce = 0%
Ended Job = job_local1327814845_0003
MapReduce Jobs Launched:
Stage-Stage-3: HDFS Read: 1542 HDFS Write: 546 SUCCESS
Total MapReduce CPU Time Spent: 0 msec
OK
1      John      Manager 50000.0 Business Management
2      Kelly     Sr. Manager 60000.0 Business Management
3      Harry     Software Engineer 45000.0 Development
4      Kedrik    Software Engineer 30000.0 Development
5      Tom       Test Engineer 20000.0 Testing
6      Ubern     Test Engineer 20000.0 Testing
7      Rohan     HR Manager 25000.0 HR
8      Joana     Software Engineer 23000.0 Development
Time taken: 51.043 seconds, Fetched: 8 row(s)
```